Hadoop Lab

Server Credentials:

master: Public Ip: 216.121.71.3, root password: Kr1sh@yPr1sh@4ggAnumeh@@*#&%9
slave1: Public Ip: 216.121.71.4, root password: Kr1sh@yPr1sh@4ggAnumeh@@*#&%9
slave2: Public Ip: 216.121.71.5, root password: Kr1sh@yPr1sh@4ggAnumeh@@*#&%9
slave3: Public Ip: 216.121.71.6, root password: Kr1sh@yPr1sh@4ggAnumeh@@*#&%9
slave4: Public Ip: 216.121.71.7, root password: Kr1sh@yPr1sh@4ggAnumeh@@*#&%9

Cloudera manager wep app credentials

admin, S0n1Gu1t@@r4gg0907

Steps for Setting up Cluster

Step1: Created 1 Extra Large 8GB, 8 cores server and 4 4 GB 4 cores server.
Step2: Changed the root password of all the servers.

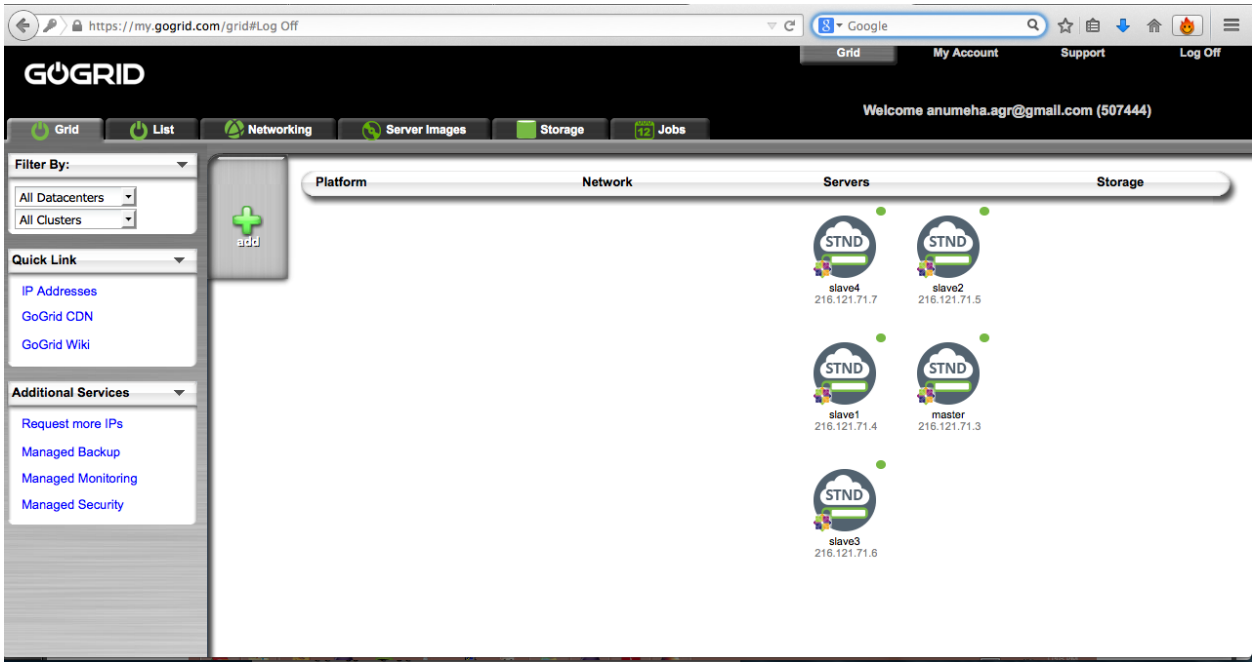

Steps for Setting up hostname configuration.

Step1: Changed the hostname for each server using hostname <hostname>,

ex. hostname master.example.com

Step2: Change the value of HOSTNAME in /etc/sysconfig/network to the <hostname>

ex. HOSTNAME=master.example.com

Step3: Restart the service network using the command

service network restart.

Step4: Open /etc/sysconfig/network-scripts/ifcfg-eth1 and make change BOOTPROTO to static and add two more values IPADDR and NETMASK. IPADDR is the private ip of the server.

```
DEVICE=eth1
BOOTPROTO=static
ONBOOT=yes
IPADDR=10.102.110.3
NETMASK=255.255.255.0
```

Step5: edit etc/hosts file and add the hostnames of all the servers. Delete the second line and add the below lines. here IP address is the private ip of the hosts.

```
10.102.110.3 master.example.com master
10.102.110.4 slave1.example.com slave1
10.102.110.5 slave2.example.com slave2
10.102.110.6 slave3.example.com slave3
10.102.110.7 slave4.example.com slave4
```

Step6: Now run the below command

ifup eth1

Step7: Next run the below command

service iptables stop

Step8: Now run the command below command

chkconfig iptables off

Step9: Make the above hostname configuration changes for all the servers

Steps For Running the cloudera manager installer.

Step1: download the latest cloudera manager installer version on the local machine.

Step2: copy the installer to the master server using SCP

scp cloudera-manager-installer.bin root@216.121.71.3:/root

Step3: Change the mode to executable of the cloudera manager installer

chmod +x cloudera-manager-installer.bin

Step4: Execute the installer.

./cloudera-manager-installer.bin

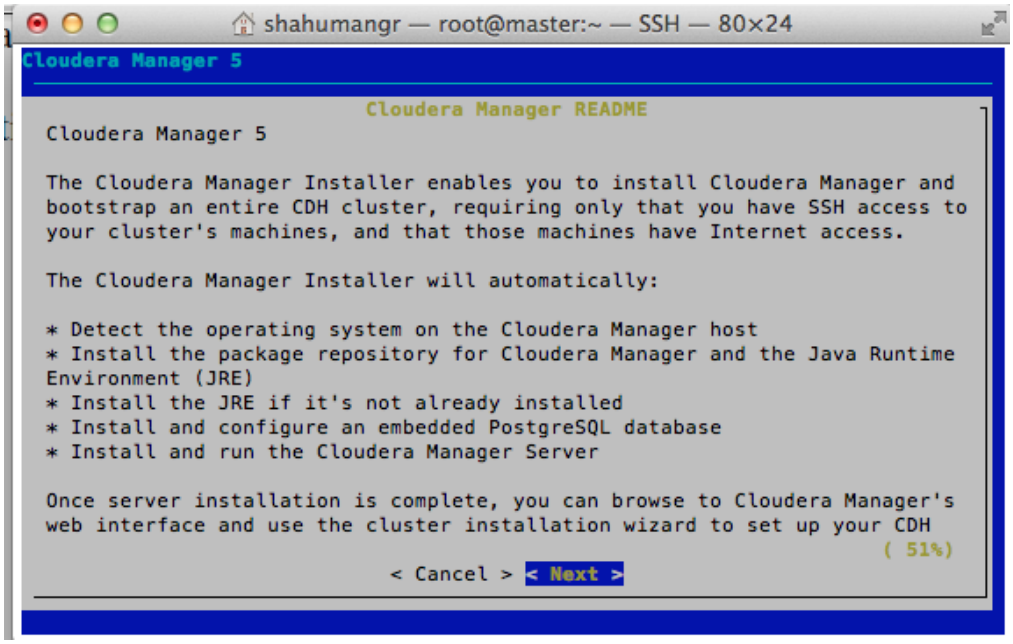

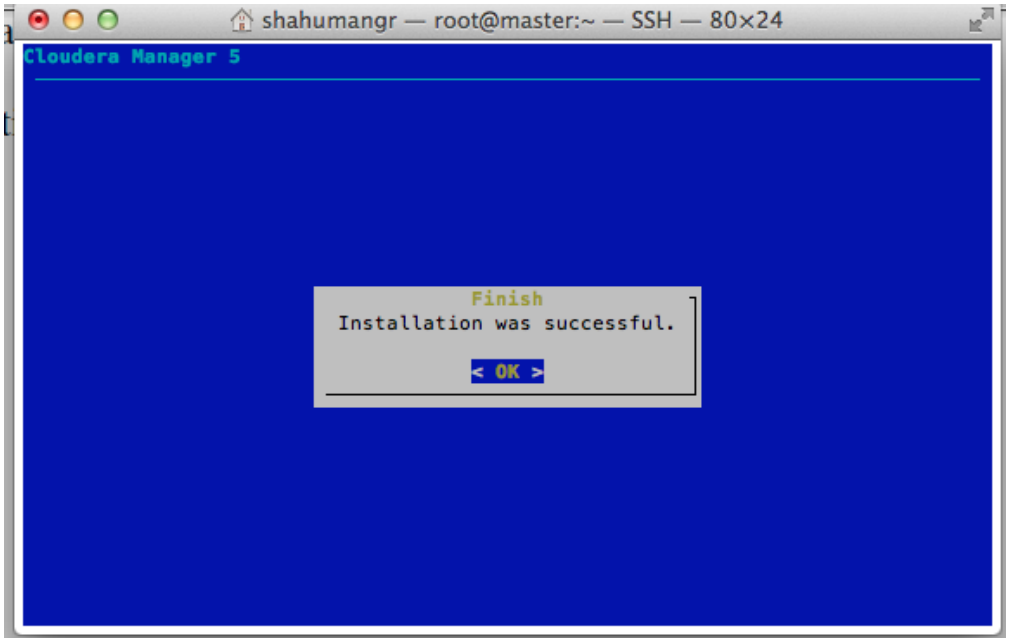

Steps For Installing CDH using cloudera manager web interface.

Step1: Login to cloudera manager web interface using admin. admin

Step2: Select the cloudera manager Data Hub Edition Trial.

Step3: Next specify hosts for your CDH cluster. Specify the private ips of the servers seperated by comma. Press continue

10.102.110.3, 10.102.110.4, 10.102.110.5, 10.102.110.6, 10.102.110.7

Step4: Next choose CDH version. I selected latest CDH 4 as it comes by default with mapreduce version 1 and CDH5 comes with mapreduce version 2 by default and mapreduce version 1 is deprecated in CDH5 but has been added for backward compatibility. Press continue

Step5: Next provide cluster information and select Login To All Hosts As: root and provide the root password for the cluster and the SSH Port which is 22 by default until and unless you changed the default SSH port. Press Continue

Once the cluster installation is finished press Continue.

Step 6: On Next page you can inspect the hosts for correctness.

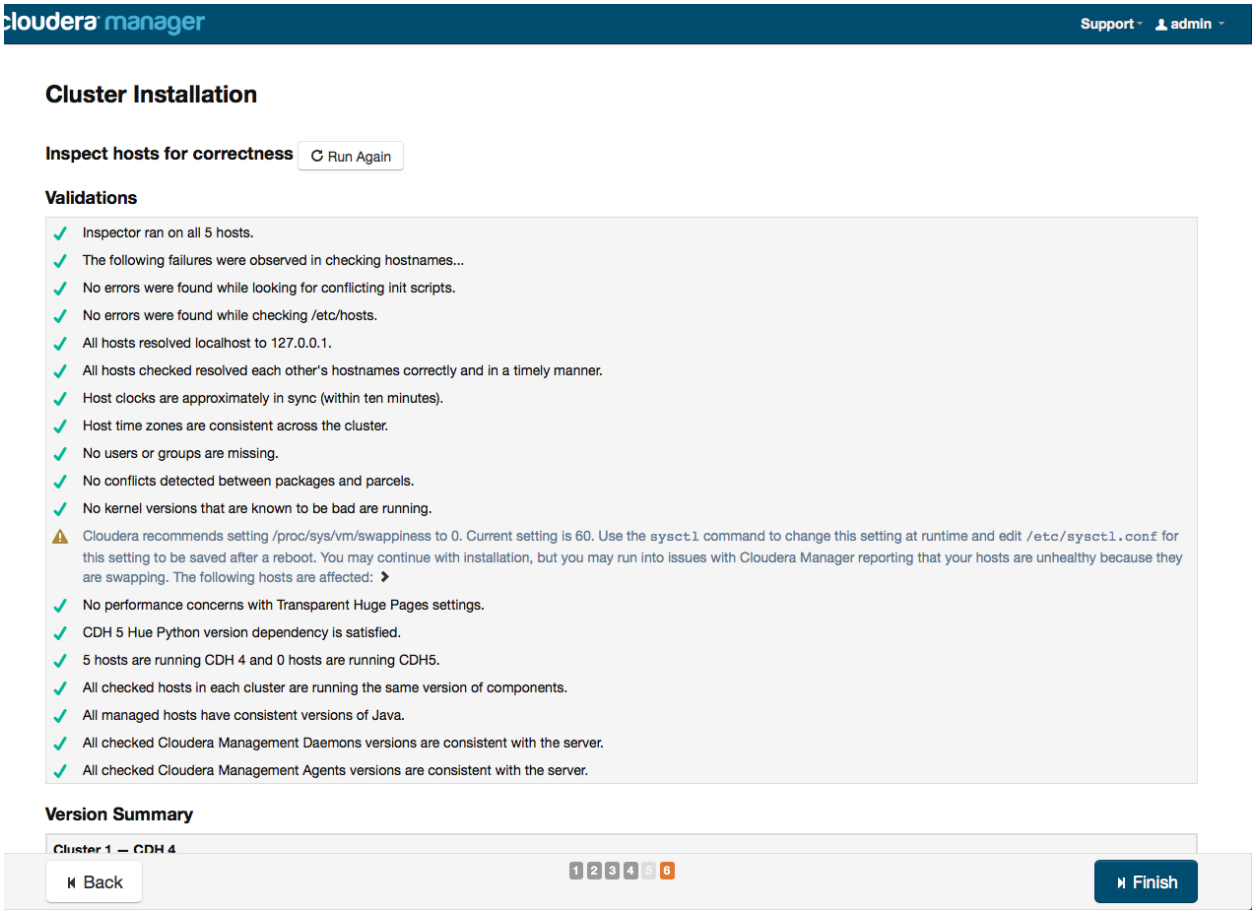

If you get a warning related to /proc/sys/vm/swappiness then set the value of swappiness accordingly using the below command

sysctl -w vm.swappiness=0.

change the value of swappiness on all the hosts.

After changing the swappiness value Click on Run Again and if you see no warning then you are good.

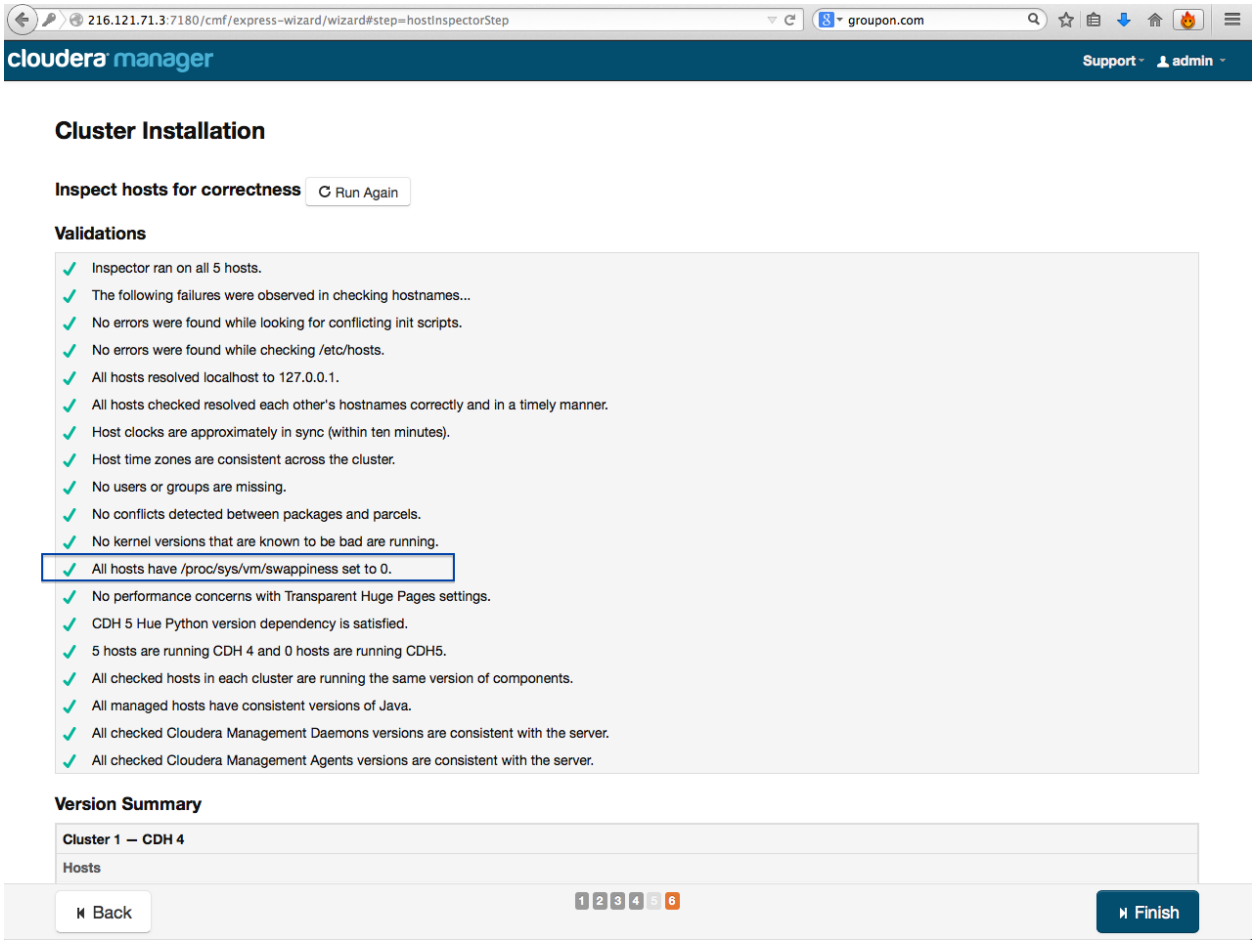

Step7: Choose the CDH4 services that you want to install. I selected All Services. On CDH5 if you want to select mapreduce instead of mapreduce 2 then select the Custom Services Option and then from the Custom Services select mapreduce.
Press Continue.

Step8: Inspect the Cluster Setup.

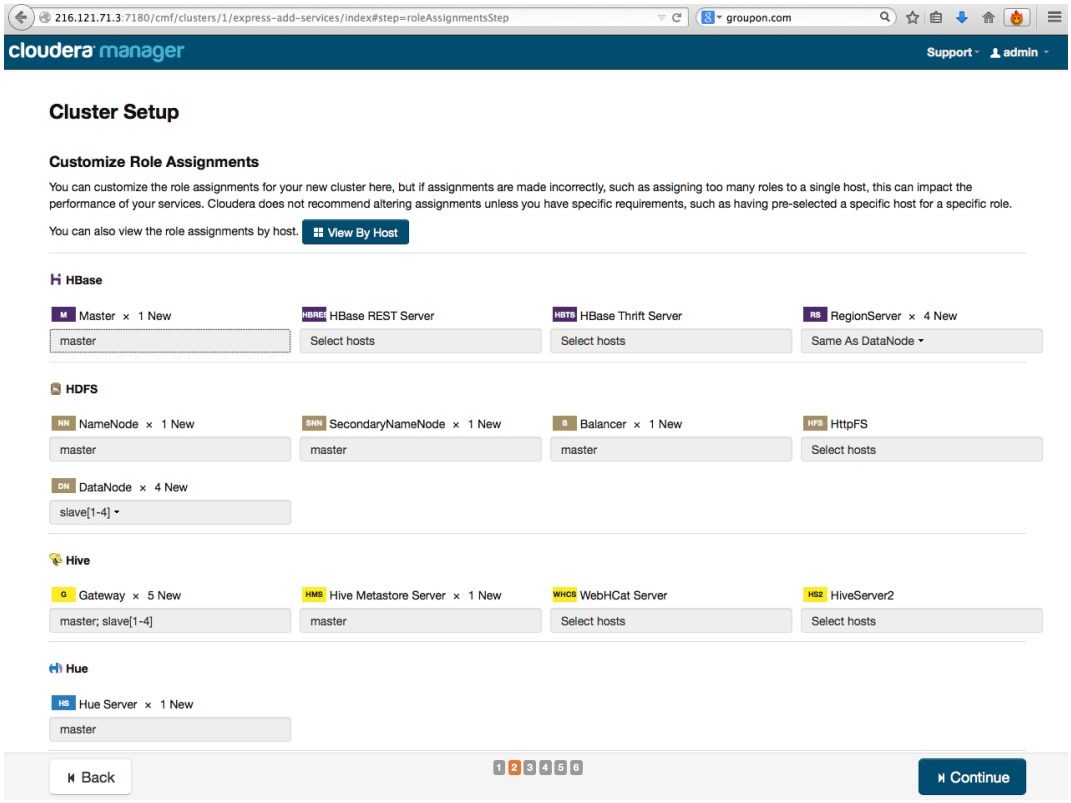

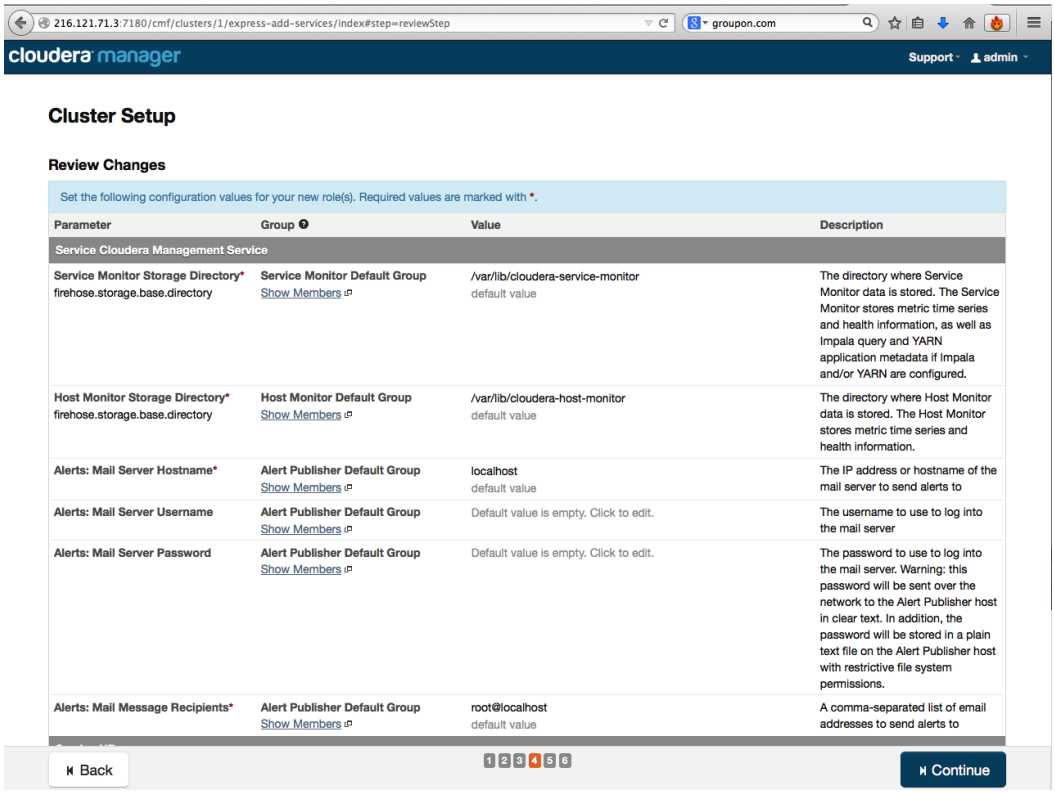

Step 9: Verify the first run of all the services. Once all the services are running Then Click on Continue.

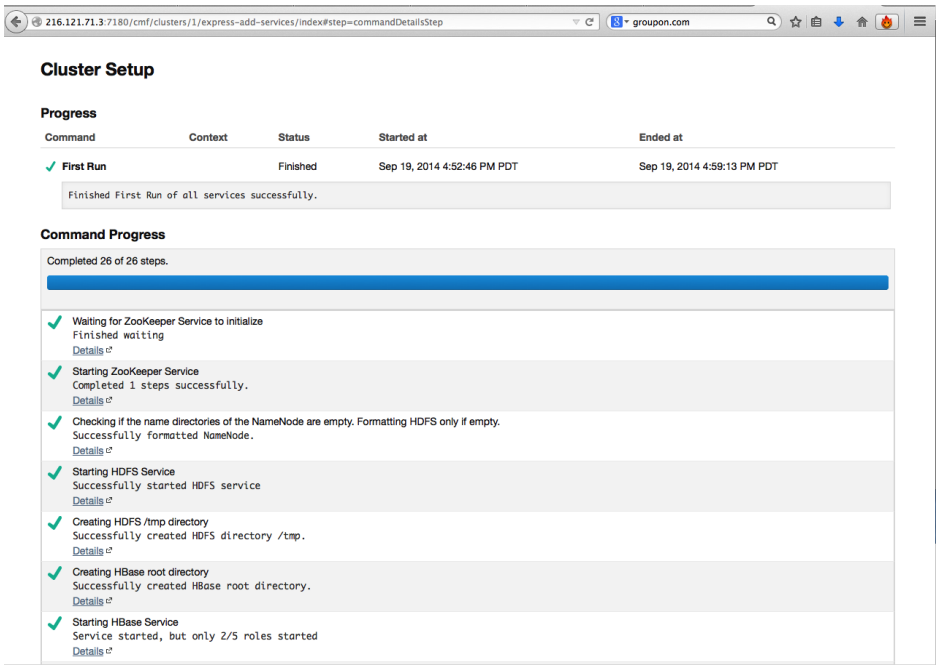

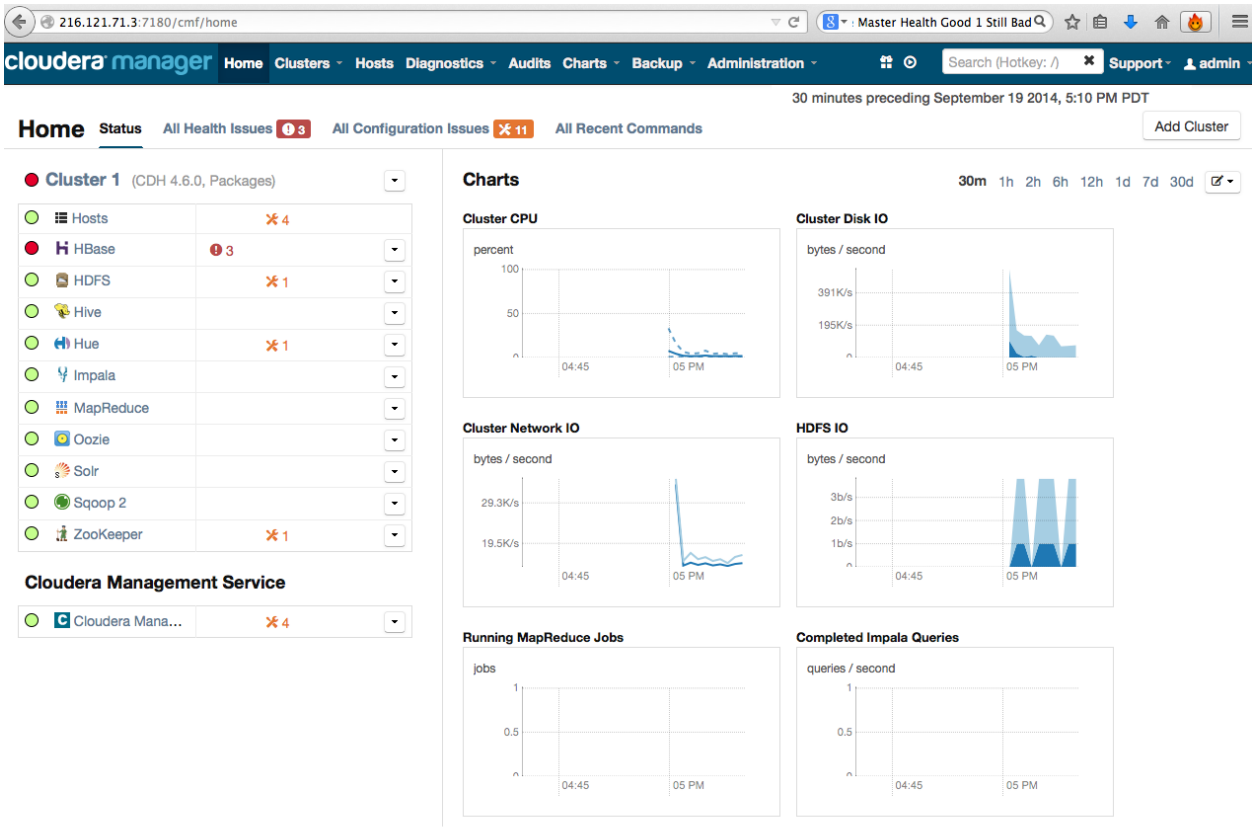

Step 10: Resolving HBase critical health issue
You might see some critical health issue for HBase.
The reason is that the slave servers are out of sync with the master servers.
We need to set the slave servers date according to master server and try to keep the difference less than 10-15 seconds.
set date using below command for all the slaves

```
date --set="Fri Sep 19 18:03:55"
```

Now the HBse issue is gone.

Output Screen shot

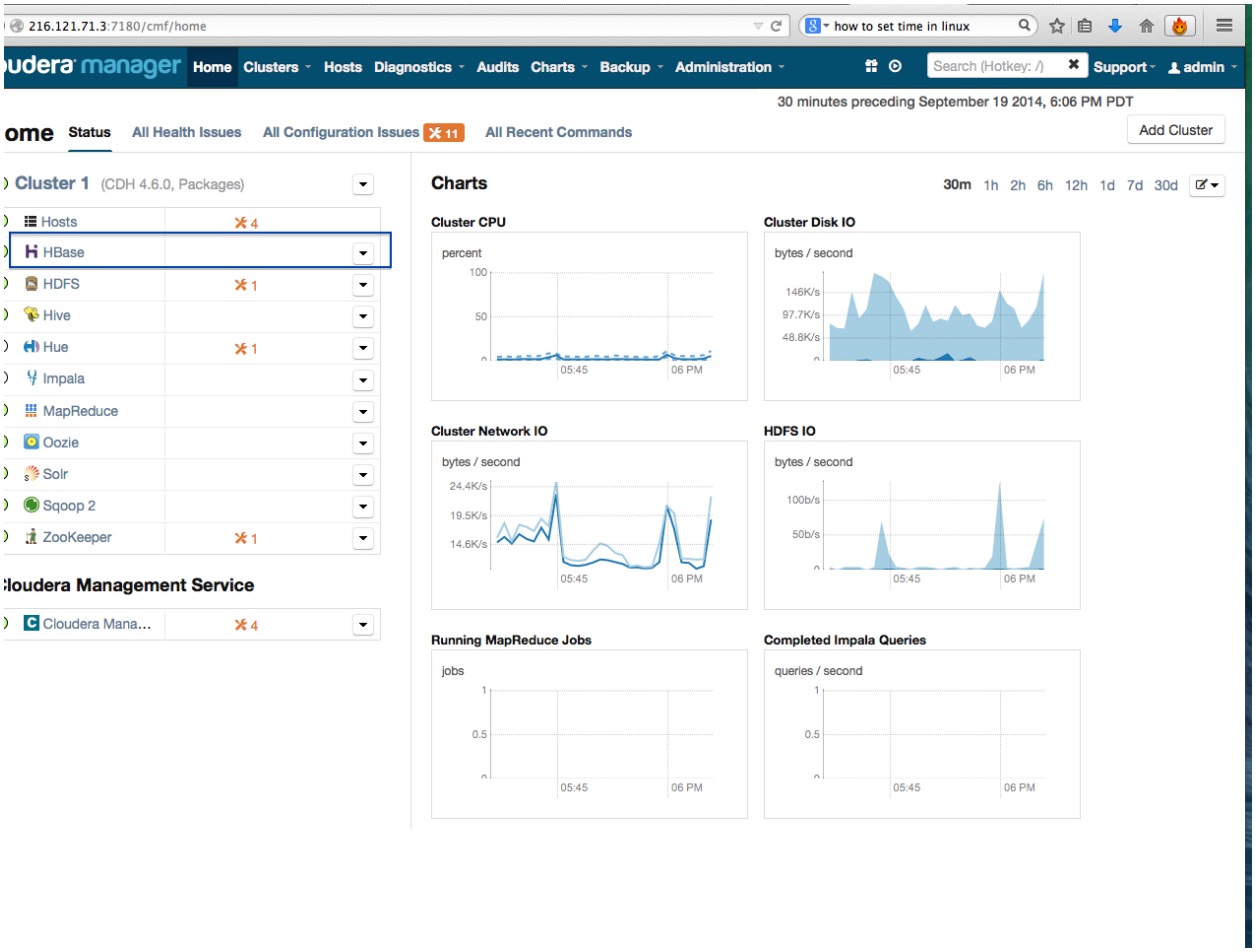

Steps For Running the Word Count Program

Step 1: Create the input directory /user/cloudera/wordcount/input in HDFS by running following commands

```
sudo su hdfs
hadoop fs -mkdir /user/cloudera
hadoop fs -chown cloudera /user/cloudera
exit
hadoop fs -mkdir /user/cloudera/wordcount /user/cloudera/wordcount/input
```

The above command will create a directory named input directory with in hdfs under /user/cloudera/wordcount.

Step 2: Now create some text files as sample files.

```
echo "Hello Hadoop Goodbye Hadoop" > file4
echo "Hello Hadoop Hello Hadoop" > file5
```

Step 3: Now Compile WordCount.java

```
mkdir wordcount_classes
javac -classpath `hadoop classpath`:. WordCount.java
```

Here hadoop classpath is the classpath for the package installation which is equivalent to /usr/lib/hadoop/*:/usr/lib/hadoop/client-0.20/*

Step 4: create the JAR of WordCount .java

```
jar -cvf wordcount.jar -C wordcount_classes/ .
```

output screen shot

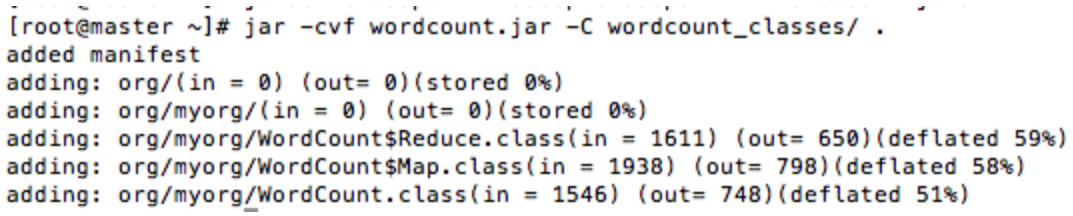

Step 5: Run the WordCount application.

```
hadoop jar wor  dcount.jar org.myorg.WordCount /user/cloudera/wordcount/input
/user/cloudera/wordcount/output31
```

output screen shot

```
[root@master ~]# hadoop jar wordcount.jar org.myorg.WordCount /user/cloudera/wordcount/input /user/cloudera/wordcount/output31
14/09/22 16:31:13 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments. Applications should implement Tool fo
14/09/22 16:31:14 INFO mapred.FileInputFormat: Total input paths to process : 2
14/09/22 16:31:16 INFO mapred.JobClient: Running job: job_201409210316_0011
14/09/22 16:31:17 INFO mapred.JobClient:  map 0% reduce 0%
14/09/22 16:31:25 INFO mapred.JobClient:  map 50% reduce 0%
14/09/22 16:31:26 INFO mapred.JobClient:  map 100% reduce 0%
14/09/22 16:31:32 INFO mapred.JobClient:  map 100% reduce 100%
14/09/22 16:31:34 INFO mapred.JobClient: Job complete: job_201409210316_0011
14/09/22 16:31:34 INFO mapred.JobClient: Counters: 33
14/09/22 16:31:34 INFO mapred.JobClient:   File System Counters
14/09/22 16:31:34 INFO mapred.JobClient:     FILE: Number of bytes read=96
14/09/22 16:31:34 INFO mapred.JobClient:     FILE: Number of bytes written=641122
14/09/22 16:31:34 INFO mapred.JobClient:     FILE: Number of read operations=0
14/09/22 16:31:34 INFO mapred.JobClient:     FILE: Number of large read operations=0
14/09/22 16:31:34 INFO mapred.JobClient:     FILE: Number of write operations=0
14/09/22 16:31:34 INFO mapred.JobClient:     HDFS: Number of bytes read=292
14/09/22 16:31:34 INFO mapred.JobClient:     HDFS: Number of bytes written=27
14/09/22 16:31:34 INFO mapred.JobClient:     HDFS: Number of read operations=6
14/09/22 16:31:34 INFO mapred.JobClient:     HDFS: Number of large read operations=0
14/09/22 16:31:34 INFO mapred.JobClient:     HDFS: Number of write operations=4
14/09/22 16:31:34 INFO mapred.JobClient:   Job Counters
14/09/22 16:31:34 INFO mapred.JobClient:     Launched map tasks=2
14/09/22 16:31:34 INFO mapred.JobClient:     Launched reduce tasks=2
14/09/22 16:31:34 INFO mapred.JobClient:     Data-local map tasks=2
14/09/22 16:31:34 INFO mapred.JobClient:     Total time spent by all maps in occupied slots (ms)=10080
14/09/22 16:31:34 INFO mapred.JobClient:     Total time spent by all reduces in occupied slots (ms)=8849
14/09/22 16:31:34 INFO mapred.JobClient:     Total time spent by all maps waiting after reserving slots (ms)=0
14/09/22 16:31:34 INFO mapred.JobClient:     Total time spent by all reduces waiting after reserving slots (ms)=0
14/09/22 16:31:34 INFO mapred.JobClient:   Map-Reduce Framework
14/09/22 16:31:34 INFO mapred.JobClient:     Map input records=2
14/09/22 16:31:34 INFO mapred.JobClient:     Map output records=8
14/09/22 16:31:34 INFO mapred.JobClient:     Map output bytes=86
14/09/22 16:31:34 INFO mapred.JobClient:     Input split bytes=238
14/09/22 16:31:34 INFO mapred.JobClient:     Combine input records=8
14/09/22 16:31:34 INFO mapred.JobClient:     Combine output records=5
14/09/22 16:31:34 INFO mapred.JobClient:     Reduce input groups=3
14/09/22 16:31:34 INFO mapred.JobClient:     Reduce shuffle bytes=128
14/09/22 16:31:34 INFO mapred.JobClient:     Reduce input records=5
14/09/22 16:31:34 INFO mapred.JobClient:     Reduce output records=3
14/09/22 16:31:34 INFO mapred.JobClient:     Spilled Records=10
14/09/22 16:31:34 INFO mapred.JobClient:     CPU time spent (ms)=3000
14/09/22 16:31:34 INFO mapred.JobClient:     Physical memory (bytes) snapshot=1094377472
14/09/22 16:31:34 INFO mapred.JobClient:     Virtual memory (bytes) snapshot=6724431872
14/09/22 16:31:34 INFO mapred.JobClient:     Total committed heap usage (bytes)=940048384
14/09/22 16:31:34 INFO mapred.JobClient:   org.apache.hadoop.mapreduce.lib.input.FileInputFormatCounter
14/09/22 16:31:34 INFO mapred.JobClient:     BYTES_READ=54
```

Step 6: Now check the output in the output directory using the below command

hadoop fs -cat /user/cloudera/wordcount/output31/part-00000

```
[root@master ~]#
[root@master ~]# hadoop fs -cat /user/cloudera/wordcount/output31/part-00000
Hadoop  4
goodbye 1
[root@master ~]# hadoop fs -cat /user/cloudera/wordcount/output31/part-00001
Hello   3
[root@master ~]#
```

Steps for running Terasort Program

Step 1: Generate the input using teragen

/usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar teragen 10 1500 /user/cloudera/terasort1/input

Output Screenshot

```
[root@master ~]# hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar teragen 10  1500 /user/cloudera/terasort/input
14/09/22 20:18:43 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments. Applications should implement Tool for the same.
Generating 10 using 2 maps with step of 5
14/09/22 20:18:45 INFO mapred.JobClient: Running job: job_201409210316_0013
14/09/22 20:18:46 INFO mapred.JobClient:  map 0% reduce 0%
14/09/22 20:18:54 INFO mapred.JobClient:  map 100% reduce 0%
14/09/22 20:18:57 INFO mapred.JobClient: Job complete: job_201409210316_0013
14/09/22 20:18:57 INFO mapred.JobClient: Counters: 24
14/09/22 20:18:57 INFO mapred.JobClient:   File System Counters
14/09/22 20:18:57 INFO mapred.JobClient:     FILE: Number of bytes read=0
14/09/22 20:18:57 INFO mapred.JobClient:     FILE: Number of bytes written=319038
14/09/22 20:18:57 INFO mapred.JobClient:     FILE: Number of read operations=0
14/09/22 20:18:57 INFO mapred.JobClient:     FILE: Number of large read operations=0
14/09/22 20:18:57 INFO mapred.JobClient:     FILE: Number of write operations=0
14/09/22 20:18:57 INFO mapred.JobClient:     HDFS: Number of bytes read=158
14/09/22 20:18:57 INFO mapred.JobClient:     HDFS: Number of bytes written=1000
14/09/22 20:18:57 INFO mapred.JobClient:     HDFS: Number of read operations=2
14/09/22 20:18:57 INFO mapred.JobClient:     HDFS: Number of large read operations=0
14/09/22 20:18:57 INFO mapred.JobClient:     HDFS: Number of write operations=2
14/09/22 20:18:57 INFO mapred.JobClient:   Job Counters
14/09/22 20:18:57 INFO mapred.JobClient:     Launched map tasks=2
14/09/22 20:18:57 INFO mapred.JobClient:     Total time spent by all maps in occupied slots (ms)=12491
14/09/22 20:18:57 INFO mapred.JobClient:     Total time spent by all reduces in occupied slots (ms)=0
14/09/22 20:18:57 INFO mapred.JobClient:     Total time spent by all maps waiting after reserving slots (ms)=0
14/09/22 20:18:57 INFO mapred.JobClient:     Total time spent by all reduces waiting after reserving slots (ms)=0
14/09/22 20:18:57 INFO mapred.JobClient:   Map-Reduce Framework
14/09/22 20:18:57 INFO mapred.JobClient:     Map input records=10
14/09/22 20:18:57 INFO mapred.JobClient:     Map output records=10
14/09/22 20:18:57 INFO mapred.JobClient:     Input split bytes=158
14/09/22 20:18:57 INFO mapred.JobClient:     Spilled Records=0
14/09/22 20:18:57 INFO mapred.JobClient:     CPU time spent (ms)=670
14/09/22 20:18:57 INFO mapred.JobClient:     Physical memory (bytes) snapshot=246104064
14/09/22 20:18:57 INFO mapred.JobClient:     Virtual memory (bytes) snapshot=3336183808
14/09/22 20:18:57 INFO mapred.JobClient:     Total committed heap usage (bytes)=214958080
14/09/22 20:18:57 INFO mapred.JobClient:   org.apache.hadoop.mapreduce.lib.input.FileInputFormatCounter
14/09/22 20:18:57 INFO mapred.JobClient:     BYTES_READ=10
```

Step 2: Next run Terasort to conduct the sorting

hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar terasort /user/root/1500 /user/cloudera/terasort/output

Output Screen Shot

```
[root@master ~]# hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar terasort /user/root/1500 /user/cloudera/terasort/output
14/09/22 20:29:16 INFO terasort.TeraSort: starting
14/09/22 20:29:17 INFO mapred.FileInputFormat: Total input paths to process : 2
14/09/22 20:29:17 INFO zlib.ZlibFactory: Successfully loaded & initialized native-zlib library
14/09/22 20:29:17 INFO compress.CodecPool: Got brand-new compressor [.deflate]
Making 2 from 10 records
Step size is 5.0
14/09/22 20:29:18 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments. Applications should implement Tool for the same.
14/09/22 20:29:21 INFO mapred.JobClient: Running job: job_201409210316_0016
14/09/22 20:29:22 INFO mapred.JobClient:  map 0% reduce 0%
14/09/22 20:29:30 INFO mapred.JobClient:  map 50% reduce 0%
14/09/22 20:29:31 INFO mapred.JobClient:  map 100% reduce 0%
14/09/22 20:29:37 INFO mapred.JobClient:  map 100% reduce 100%
14/09/22 20:29:39 INFO mapred.JobClient: Job complete: job_201409210316_0016
14/09/22 20:29:39 INFO mapred.JobClient: Counters: 33
14/09/22 20:29:39 INFO mapred.JobClient:   File System Counters
14/09/22 20:29:39 INFO mapred.JobClient:     FILE: Number of bytes read=793
14/09/22 20:29:39 INFO mapred.JobClient:     FILE: Number of bytes written=651886
14/09/22 20:29:39 INFO mapred.JobClient:     FILE: Number of read operations=0
14/09/22 20:29:39 INFO mapred.JobClient:     FILE: Number of large read operations=0
14/09/22 20:29:39 INFO mapred.JobClient:     FILE: Number of write operations=0
14/09/22 20:29:39 INFO mapred.JobClient:     HDFS: Number of bytes read=1218
14/09/22 20:29:39 INFO mapred.JobClient:     HDFS: Number of bytes written=1000
14/09/22 20:29:39 INFO mapred.JobClient:     HDFS: Number of read operations=4
14/09/22 20:29:39 INFO mapred.JobClient:     HDFS: Number of large read operations=0
14/09/22 20:29:39 INFO mapred.JobClient:     HDFS: Number of write operations=2
14/09/22 20:29:39 INFO mapred.JobClient:   Job Counters
14/09/22 20:29:39 INFO mapred.JobClient:     Launched map tasks=2
14/09/22 20:29:39 INFO mapred.JobClient:     Launched reduce tasks=2
14/09/22 20:29:39 INFO mapred.JobClient:     Data-local map tasks=2
14/09/22 20:29:39 INFO mapred.JobClient:     Total time spent by all maps in occupied slots (ms)=10234
14/09/22 20:29:39 INFO mapred.JobClient:     Total time spent by all reduces in occupied slots (ms)=8389
14/09/22 20:29:39 INFO mapred.JobClient:     Total time spent by all maps waiting after reserving slots (ms)=0
14/09/22 20:29:39 INFO mapred.JobClient:     Total time spent by all reduces waiting after reserving slots (ms)=0
14/09/22 20:29:39 INFO mapred.JobClient:   Map-Reduce Framework
14/09/22 20:29:39 INFO mapred.JobClient:     Map input records=10
14/09/22 20:29:39 INFO mapred.JobClient:     Map output records=10
14/09/22 20:29:39 INFO mapred.JobClient:     Map output bytes=1000
14/09/22 20:29:39 INFO mapred.JobClient:     Input split bytes=218
14/09/22 20:29:39 INFO mapred.JobClient:     Combine input records=0
14/09/22 20:29:39 INFO mapred.JobClient:     Combine output records=0
14/09/22 20:29:39 INFO mapred.JobClient:     Reduce input groups=10
14/09/22 20:29:39 INFO mapred.JobClient:     Reduce shuffle bytes=547
14/09/22 20:29:39 INFO mapred.JobClient:     Reduce input records=10
14/09/22 20:29:39 INFO mapred.JobClient:     Reduce output records=10
14/09/22 20:29:39 INFO mapred.JobClient:     Spilled Records=20
14/09/22 20:29:39 INFO mapred.JobClient:     CPU time spent (ms)=2650
```

```
14/09/22 20:29:39 INFO mapred.JobClient:     CPU time spent (ms)=2650
14/09/22 20:29:39 INFO mapred.JobClient:     Physical memory (bytes) snapshot=1111457792
14/09/22 20:29:39 INFO mapred.JobClient:     Virtual memory (bytes) snapshot=6739128320
14/09/22 20:29:39 INFO mapred.JobClient:     Total committed heap usage (bytes)=940048384
14/09/22 20:29:39 INFO mapred.JobClient:   org.apache.hadoop.mapreduce.lib.input.FileInputFormatCounter
14/09/22 20:29:39 INFO mapred.JobClient:     BYTES_READ=1000
14/09/22 20:29:39 INFO terasort.TeraSort: done
```

Note : You might run into java versioning problem Steps for handling java versioning problem

Step 1: Make sure that your servers are using java 1.7.x
Step2: set JAVA_HOME and point it to java version 1.7.x on all your servers.

Step 3: If you are still getting versioning issue then change the JAVA_HOME values in following two files and point it to java 1.7.x.

Step 4: Run the above process for all the servers
Step 5: Restart the server